

Red Wine Quality Prediction Analysis

March 30, 2024

1 Wine Quality Prediction Analysis

The two datasets are related to red and white variants of the Portuguese “Vinho Verde” wine. These datasets can be viewed as **classification or regression** tasks. The classes are ordered and not balanced (e.g. there are much more normal wines than excellent or poor ones). Outlier detection algorithms could be used to detect the few excellent or poor wines. Also, we are not sure if all input variables are relevant. So it could be interesting to test feature selection methods. Two datasets were combined and few values were randomly removed

1.1 Attribute Information:

Input variables (based on physicochemical tests):

1. **fixed acidity**
2. **volatile acidity**
3. **citric acid**
4. **residual sugar**
5. **chlorides**
6. **free sulfur dioxide**
7. **total sulfur dioxide**
8. **density**
9. **pH**
10. **sulphates**
11. **alcohol**

Output variable (based on sensory data):

12. **quality** (score between 0 and 10)

1.2 Import Modules

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')
```

1.3 Loading The Dataset

```
[2]: df = pd.read_csv('winequality.csv')
df.head()
```

```
[2]:
```

	type	fixed acidity	volatile acidity	citric acid	residual sugar	\
0	white	7.0	0.27	0.36	20.7	
1	white	6.3	0.30	0.34	1.6	
2	white	8.1	0.28	0.40	6.9	
3	white	7.2	0.23	0.32	8.5	
4	white	7.2	0.23	0.32	8.5	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	\
0	0.045	45.0	170.0	1.0010	3.00	
1	0.049	14.0	132.0	0.9940	3.30	
2	0.050	30.0	97.0	0.9951	3.26	
3	0.058	47.0	186.0	0.9956	3.19	
4	0.058	47.0	186.0	0.9956	3.19	

	sulphates	alcohol	quality
0	0.45	8.8	6
1	0.49	9.5	6
2	0.44	10.1	6
3	0.40	9.9	6
4	0.40	9.9	6

```
[3]: # statistical information
df.describe()
```

```
[3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	6487.000000	6489.000000	6494.000000	6495.000000	
mean	7.216579	0.339691	0.318722	5.444326	
std	1.296750	0.164649	0.145265	4.758125	
min	3.800000	0.080000	0.000000	0.600000	
25%	6.400000	0.230000	0.250000	1.800000	
50%	7.000000	0.290000	0.310000	3.000000	
75%	7.700000	0.400000	0.390000	8.100000	
max	15.900000	1.580000	1.660000	65.800000	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	\
count	6495.000000	6497.000000	6497.000000	6497.000000	
mean	0.056042	30.525319	115.744574	0.994697	
std	0.035036	17.749400	56.521855	0.002999	
min	0.009000	1.000000	6.000000	0.987110	
25%	0.038000	17.000000	77.000000	0.992340	
50%	0.047000	29.000000	118.000000	0.994890	
75%	0.065000	41.000000	156.000000	0.996990	
max	0.611000	289.000000	440.000000	1.038980	

	pH	sulphates	alcohol	quality
count	6488.000000	6493.000000	6497.000000	6497.000000
mean	3.218395	0.531215	10.491801	5.818378
std	0.160748	0.148814	1.192712	0.873255
min	2.720000	0.220000	8.000000	3.000000
25%	3.110000	0.430000	9.500000	5.000000
50%	3.210000	0.510000	10.300000	6.000000
75%	3.320000	0.600000	11.300000	6.000000
max	4.010000	2.000000	14.900000	9.000000

```
[4]: # Datatype Information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   type                   6497 non-null   object
1   fixed acidity          6487 non-null   float64
2   volatile acidity       6489 non-null   float64
3   citric acid            6494 non-null   float64
4   residual sugar         6495 non-null   float64
5   chlorides              6495 non-null   float64
6   free sulfur dioxide    6497 non-null   float64
7   total sulfur dioxide   6497 non-null   float64
8   density                6497 non-null   float64
9   pH                     6488 non-null   float64
10  sulphates              6493 non-null   float64
11  alcohol                6497 non-null   float64
12  quality                6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

```
[5]: # Cheak for NULL values
df.isnull().sum()
```

```
[5]: type                0
fixed acidity          10
volatile acidity       8
citric acid            3
residual sugar         2
chlorides              2
free sulfur dioxide    0
total sulfur dioxide   0
density               0
pH                    9
```

```
sulphates          4
alcohol            0
quality            0
dtype: int64
```

```
[6]: # fill the missing values
for col, values in df.items():
    if col != 'type':
        df[col] = df[col].fillna(df[col].mean())
```

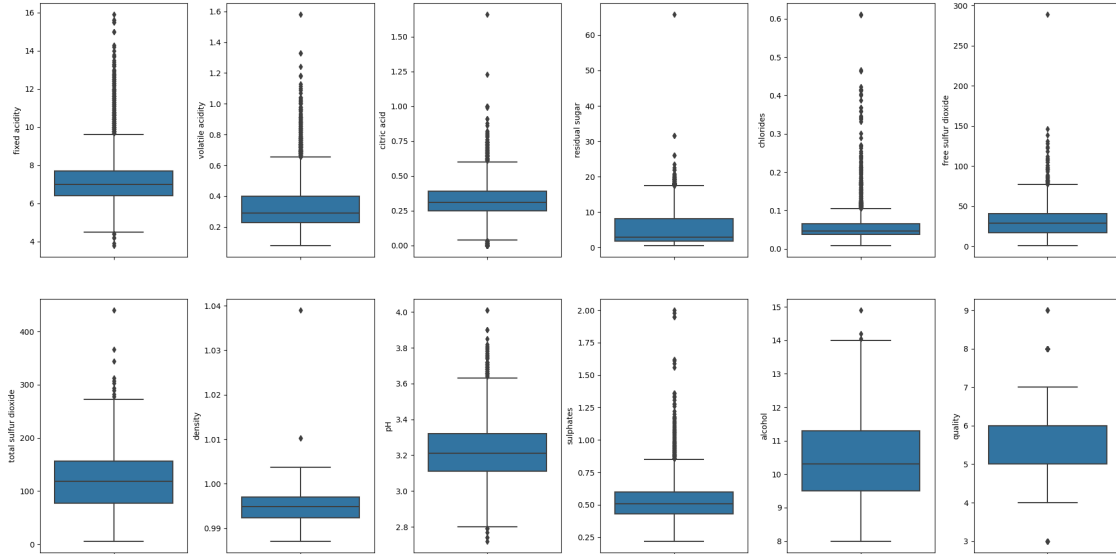
```
[7]: # Cheak for NULL values ( after fill missing values)
df.isnull().sum()
```

```
[7]: type          0
fixed acidity     0
volatile acidity  0
citric acid       0
residual sugar    0
chlorides         0
free sulfur dioxide 0
total sulfur dioxide 0
density          0
pH               0
sulphates        0
alcohol          0
quality          0
dtype: int64
```

1.4 Expoloratory Data Analysis

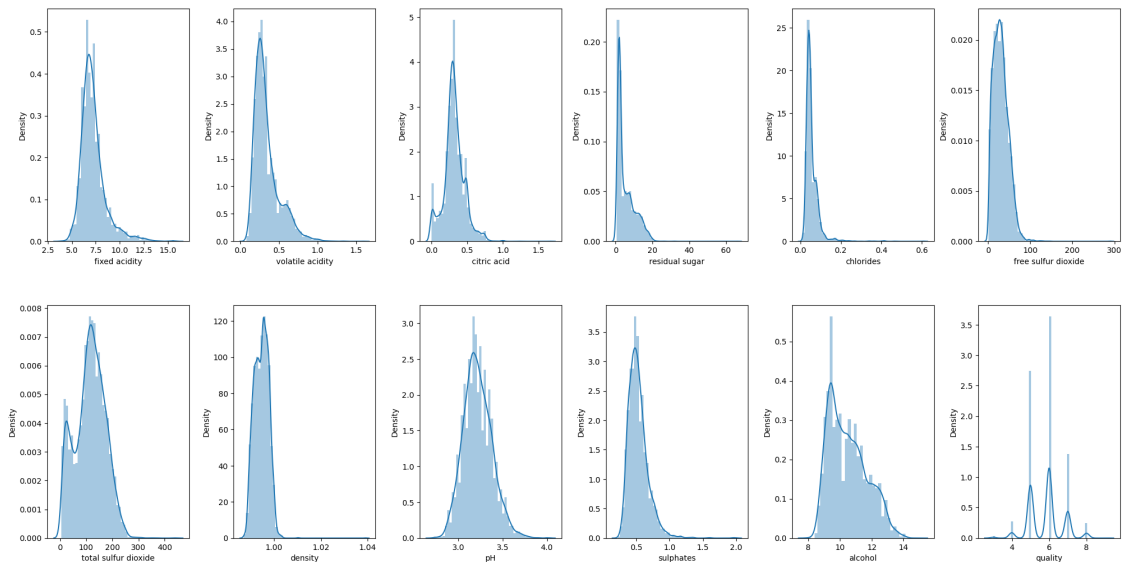
```
[8]: # create box plots
fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col, value in df.items():
    if col != 'type':
        sns.boxplot(y=col, data=df, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



```
[9]: # create dist plot
fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()


for col, value in df.items():
    if col != 'type':
        sns.distplot(value, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



```
[10]: # log transformation
df['free sulfur dioxide'] = np.log(1 + df['free sulfur dioxide'])
```

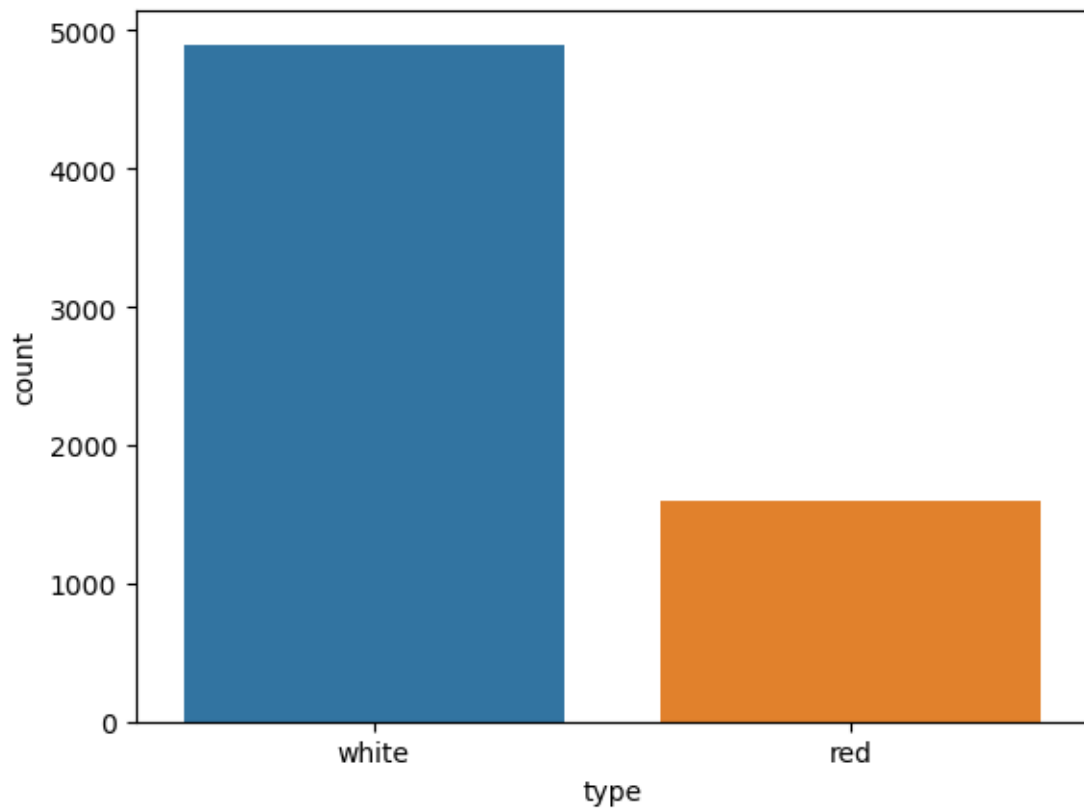
```
[11]: sns.distplot(df['free sulfur dioxide'])
```

```
[11]: <Axes: xlabel='free sulfur dioxide', ylabel='Density'>
```




```
[12]: sns.countplot(data=df, x='type')
```

```
[12]: <Axes: xlabel='type', ylabel='count'>
```



```
[13]: sns.countplot(data=df, x='quality')
```

```
[13]: <Axes: xlabel='quality', ylabel='count'>
```




1.5 Coorelation Matrix

```
[14]: # Step 1: Identify non-numeric columns
non_numeric_columns = df.select_dtypes(exclude=['float64', 'int64']).columns

# Step 2: Drop non-numeric columns
numeric_df = df.drop(columns=non_numeric_columns)

# Now, calculate the correlation matrix
corr = numeric_df.corr()

# Plot heatmap
plt.figure(figsize=(20, 10))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.show()
```

1.6 Input split

```
[15]: X = df.drop(columns=['type', 'quality'])
      y = df['quality']
```

1.7 Class Imbalancement

```
[16]: y.value_counts()
```

```
[16]: quality
6      2836
5      2138
7      1079
4       216
8       193
3        30
9         5
Name: count, dtype: int64
```

```
[17]: from imblearn.over_sampling import SMOTE
      oversample = SMOTE(k_neighbors=4)
      # transform the dataset
      X, y = oversample.fit_resample(X, y)
```

```
[18]: y.value_counts()
```

```
[18]: quality
      6      2836
      5      2836
      7      2836
      8      2836
      4      2836
      3      2836
      9      2836
      Name: count, dtype: int64
```

1.8 Model Training

```
[19]: # classify function
from sklearn.model_selection import cross_val_score, train_test_split
def classify(model, X, y):
    x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    random_state=42)
    # train the model
    model.fit(x_train, y_train)
    print("Accuracy:", model.score(x_test, y_test) * 100)

    # cross-validation
    score = cross_val_score(model, X, y, cv=5)
    print("CV Score:", np.mean(score)*100)
```

```
[20]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
classify(model, X, y)
```

Accuracy: 35.05943985492645
CV Score: 32.81264482358561

```
[21]: from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier()
classify(model, X, y)
```

Accuracy: 80.19343139230304
CV Score: 75.71034965718081

```
[22]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
classify(model, X, y)
```

Accuracy: 88.01128349788434
CV Score: 82.72220957102722

```
[23]: from sklearn.ensemble import ExtraTreesClassifier
model = ExtraTreesClassifier()
classify(model, X, y)
```

Accuracy: 88.87769494257506
CV Score: 83.65408278025761

[]:

[]: